

ICP3

Tejaswini Tatikonda – 700762235

GitHub link: <https://github.com/tejaswini22350/Assignment-3.git>

Video Link:

https://drive.google.com/file/d/1kHuDmVpyCWly1ORw7Ge3tyVT_y9ZDXtE/view?usp=share_link

1. Data Manipulation

- a. Read the provided CSV file 'data.csv'.
- b. <https://drive.google.com/drive/folders/1h8C3mLsso-R-sIOLsvoYwPLzy2fJ4IOF?usp=sharing>
- c. Show the basic statistical description about the data.
- d. Check if the data has null values.
- i. Replace the null values with the mean
- e. Select at least two columns and aggregate the data using: min, max, count, mean.
- f. Filter the dataframe to select the rows with calories values between 500 and 1000.
- g. Filter the dataframe to select the rows with calories values > 500 and pulse < 100 .
- h. Create a new "df_modified" dataframe that contains all the columns from df except for "Maxpulse".
- i. Delete the "Maxpulse" column from the main df dataframe
- j. Convert the datatype of Calories column to int datatype.
- k. Using pandas create a scatter plot for the two columns (Duration and Calories).

```

import pandas as pd
import matplotlib.pyplot as plt
df = pd.read_csv('/content/data.csv')
print("Basic Statistical Description:")
print(df.describe())
print("\nCheck for Null Values:")
print(df.isnull().sum())
df.fillna(df.mean(), inplace=True)
agg_columns = ['Calories', 'Duration']
agg_result = df[agg_columns].agg(['min', 'max', 'count', 'mean'])
print("\nAggregated Data:")
print(agg_result)
filtered_df1 = df[(df['Calories'] >= 500) & (df['Calories'] <= 1000)]
filtered_df2 = df[(df['Calories'] > 500) & (df['Pulse'] < 100)]
df_modified = df.drop(columns=['Maxpulse'])
df.drop(columns=['Maxpulse'], inplace=True)
df['Calories'] = df['Calories'].astype(int)
df.plot.scatter(x='Duration', y='Calories', title='Scatter Plot: Duration vs Calories')
plt.show()

```

```

Basic Statistical Description:

```

| | Duration | Pulse | Maxpulse | Calories |
|-------|------------|------------|------------|-------------|
| count | 169.000000 | 169.000000 | 169.000000 | 164.000000 |
| mean | 63.846154 | 107.461538 | 134.047337 | 375.790244 |
| std | 42.299949 | 14.510259 | 16.450434 | 266.379919 |
| min | 15.000000 | 80.000000 | 100.000000 | 50.300000 |
| 25% | 45.000000 | 100.000000 | 124.000000 | 250.925000 |
| 50% | 60.000000 | 105.000000 | 131.000000 | 318.600000 |
| 75% | 60.000000 | 111.000000 | 141.000000 | 387.600000 |
| max | 300.000000 | 159.000000 | 184.000000 | 1860.400000 |

```

Check for Null Values:
Duration      0
Pulse         0
Maxpulse      0
Calories      5
dtype: int64

```

```

Aggregated Data:

```

| | Calories | Duration |
|-------|-------------|------------|
| min | 50.300000 | 15.000000 |
| max | 1860.400000 | 300.000000 |
| count | 169.000000 | 169.000000 |
| mean | 375.790244 | 63.846154 |

```

Basic Statistical Description:

```

| | Duration | Pulse | Maxpulse | Calories |
|-------|------------|------------|------------|-------------|
| count | 169.000000 | 169.000000 | 169.000000 | 164.000000 |
| mean | 63.846154 | 107.461538 | 134.047337 | 375.790244 |
| std | 42.299949 | 14.510259 | 16.450434 | 266.379919 |
| min | 15.000000 | 80.000000 | 100.000000 | 50.300000 |
| 25% | 45.000000 | 100.000000 | 124.000000 | 250.925000 |
| 50% | 60.000000 | 105.000000 | 131.000000 | 318.600000 |
| 75% | 60.000000 | 111.000000 | 141.000000 | 387.600000 |
| max | 300.000000 | 159.000000 | 184.000000 | 1860.400000 |

```

Check for Null Values:
Duration      0
Pulse         0
Maxpulse      0
Calories      5
dtype: int64


```

```

Aggregated Data:

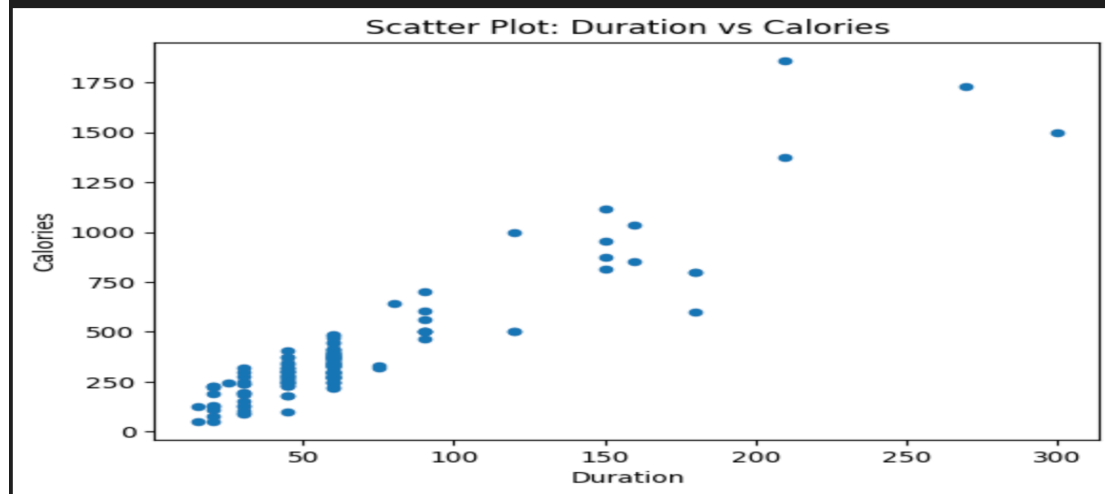
```

| | Calories | Duration |
|-------|-------------|------------|
| min | 50.300000 | 15.000000 |
| max | 1860.400000 | 300.000000 |
| count | 169.000000 | 169.000000 |
| mean | 375.790244 | 63.846154 |



Aggregated Data:

| | Calories | Duration |
|-------|-------------|------------|
| min | 50.300000 | 15.000000 |
| max | 1860.400000 | 300.000000 |
| count | 169.000000 | 169.000000 |
| mean | 375.790244 | 63.846154 |



2. Linear Regression

- Import the given "Salary_Data.csv"
- Split the data in train_test partitions, such that 1/3 of the data is reserved as test subset.
- Train and predict the model.
- Calculate the mean_squared error

e) Visualize both train and test data using scatter plot.

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
import matplotlib.pyplot as plt

d = '/Salary_Data (2) (1) (1).csv'
df = pd.read_csv(d)
print("First few rows of the dataframe:")
print(df.head())
X = df[['YearsExperience']]
y = df['Salary']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=1/3, random_state=42)
model = LinearRegression()
model.fit(X_train, y_train)
y_train_pred = model.predict(X_train)
y_test_pred = model.predict(X_test)
mse_train = mean_squared_error(y_train, y_train_pred)
mse_test = mean_squared_error(y_test, y_test_pred)
print(f"\nMean Squared Error (Train): {mse_train}")
print(f"Mean Squared Error (Test): {mse_test}")
plt.scatter(X_train, y_train, label='Train Data', color='blue')
plt.scatter(X_test, y_test, label='Test Data', color='red', marker='x')
plt.plot(X_train, y_train_pred, label='Regression Line', color='green')
plt.title('Salary vs Years of Experience')
plt.xlabel('Years of Experience')
plt.ylabel('Salary')
plt.legend()
plt.show()
```

[47]

```
... First few rows of the dataframe:
   YearsExperience  Salary
0              1.1  39343.0
1              1.3  46205.0
2              1.5  37731.0
3              2.0  43525.0
4              2.2  39891.0

Mean Squared Error (Train): 29793161.082422983
Mean Squared Error (Test): 35301898.887134895
```

